Governance, Risk & Responsible Use: A Crash Course

Foundations (Everyone), Chapter 9/10 — Revision & Viva-Prep Cheatsheet (English + Roman Urdu)
Source: agentfactory.panaversity.org/docs/governance-risk-responsible-use-crash-course

EN: 4 questions decide how and where AI does meaningful work: the Case (should this be delegated?), the Data (what can it see?), the Capability (what can it do?), the People (who's affected?).
Roman Urdu: 4 sawal decide karte AI meaningful work kahan/kaise kare: the Case (delegate honi chahiye?), the Data (kya dekh sakti hai?), the Capability (kya kar sakti hai?), the People (kaun affected hai?).

1. The Case + The Data

- **Delegation criteria:** reversibility, consequence, human-judgment need, accountability. A defined gate (who/what/when) beats 'someone will check it'.
- **3 data tiers:** green/yellow/red. Discriminator: identifiers vs. pattern. Redaction has 2 failure modes (too much removed = useless; too little = still identifying).

2. The Capability + The People

- Skills are software with no permission manifest (unlike connector scopes) -- **5 checks:** source, reach, fit, outside-content, actions. 3 outcomes: enable, escalate, decline.
- Prompt injection has 2 trust axes. For People: look at who was EXCLUDED (shortlist blind spot) -- rules first, judgment second. Escalate the question, not a verdict.

3. One Workflow, Governance Record, Incident Response

- Ayesha's logistics service-exception-report example walks all 4 questions end to end.
- **First-hour incident procedure:** stop, record, report, state facts, follow the owner. Near-misses matter too, not just failures.

4. Why Governance Drifts

- The diligence gap -- usage audits should audit the PROCESS, not the person. Friction rule: too much friction breeds shadow AI use outside policy.

Key Terms Cheatsheet

One line each — English + Roman Urdu.

Term	English	Roman Urdu
4 governance questions	The Case, the Data, the Capability, the People.	The Case, the Data, the Capability, the People.
3 data tiers	Green / yellow / red, discriminated by identifiers vs. pattern.	Green / yellow / red, identifiers vs. pattern se discriminate.
Redaction's 2 failure modes	Over-redaction (useless) and under-redaction (still identifying).	Over-redaction (useless) aur under-redaction (phir bhi identifying).
Shadow AI	Unsanctioned AI use that happens when official policy has too much friction.	Unsanctioned AI use jo tab hota jab official policy mein bohot friction ho.
Governance Record	A 1-page template documenting the 4-question decision and evidence.	1-page template jo 4-question decision + evidence document karta.

Quick Self-Check

Question	Answer
What are the 4 governance questions this chapter runs on every workflow?	The Case (should it be delegated?), the Data (what can it see?), the Capability (what can it do?), the People (who's affected?).
What's wrong with 'someone will check it' as a review plan?	It's not a defined gate -- it doesn't say who checks, what they check, or when, so it quietly fails.
Why does a usage audit look at the PROCESS, not the person?	Blaming individuals misses systemic gaps; auditing the process finds where governance actually breaks down.
Why does over-strict AI policy cause its own risk?	High friction pushes people toward unsanctioned 'shadow AI' use, which is harder to govern than sanctioned use.

Where This Leads / Sources

Leads to: Human-Agent Teams, Cowork, Designing Agent Experiences, System of Record. Full 21-Q quiz in 09-test-your-understanding.md.

Source: docs/governance-risk-responsible-use-crash-course/ (this repo), content originally via Zia Tutor AI connector.