AI Fluency: A Crash Course

Revision & Viva-Prep Cheatsheet — English + Roman Urdu

Core one-liner: Decide what AI should do (Delegation). Describe the work clearly (Description). Check what comes back (Discernment). Own what happens next (Diligence).
Roman Urdu: AI kya kare decide karo. Kaam clearly describe karo. Jo wapas aaye usay check karo. Aage jo hota hai uski responsibility lo.

1. What is AI Fluency?

EN: AI fluency means working with AI in a way that is Effective (reaches the goal), Efficient (no wasted time/tokens), Ethical (fair, open use), and Safe (protects people, privacy, security). AI access (having the tool) is not the same as AI fluency (knowing what to do with it).

Roman Urdu: AI fluency matlab AI ke sath is tarah kaam karna jo Effective (goal reach ho), Efficient (time/tokens waste na hon), Ethical (fair/open use), aur Safe (log/privacy/security protect ho) ho. AI access (tool hona) aur AI fluency (uska sahi use pata hona) alag cheezein hain.

Framework by Prof. Rick Dakan + Prof. Joseph Feller, with Anthropic — 4 human competencies: **Delegation, Description, Discernment, Diligence** (the 4Ds).

2. Three Ways to Work with AI (Modes)

Mode	One-line Definition (EN)	Roman Urdu
Automation “Do this task”	AI performs a task you specify exactly; you are the script writer. Fits clear, repeatable work.	AI specify kiya task perform karta hai; tum script writer ho. Clear/repeatable kaam ke liye fit.
Augmentation “Help me think”	Human and AI act as thinking partners across multiple turns: ask, respond, challenge, revise.	Insaan aur AI thinking partners ki tarah kaam karte hain: ask-respond-challenge-revise.
Agency “Pursue this goal for me”	AI is configured to independently perform FUTURE tasks, including FOR OTHERS. You're the director, but you're not in the room while it runs.	AI future tasks khud independently kare, others ke liye bhi — tum director ho, par room mein nahi ho.

Key distinction: in Agency the work is **FUTURE** + **FOR OTHERS** — judgment must be built in beforehand since you won't supervise each step.

3. The 4Ds Framework

Delegation = decide what AI should do • **Description** = explain what's needed • **Discernment** = judge what AI gave you • **Diligence** = use responsibly, own the outcome.

3.1 Delegation — Deciding the Human/AI Split

EN: Delegation means deciding how work divides between human and AI — designing the workflow, not just handing AI a task. The most common beginner mistake happens BEFORE the first prompt: typing without deciding what's needed, what good looks like, and which decisions must stay human.

Roman Urdu: Delegation = decide karna kaam insaan aur AI mein kaise divide ho — workflow design karna, sirf AI ko task dena nahi.

- **Problem awareness (EN):** Understanding the goal, audience, success criteria, risks, and where human judgment is essential.
  Roman Urdu: Goal, audience, success, risks, aur human judgment kahan zaroori hai — yeh samajhna.
- **Platform awareness (EN):** Knowing which AI system/tool fits this job (reasoning model, search-enabled, coding agent, etc).
  Roman Urdu: Kaunsa AI system/tool is kaam ke liye fit karta hai, jaanna.
- **Task delegation (EN):** Deciding which parts AI does, which parts you do, and why — not “can AI do this?” but “which parts, and why?”
  Roman Urdu: Decide karna kaunsa part AI kare, kaunsa tum karo, aur kyun.

Example: “Build me an invoice-chasing agent” is not enough — tone, late-days threshold, approval amount, dispute handling are business policy questions AI cannot decide for you.

3.2 Description — Giving AI What It Needs

EN: Description means giving AI the information and guidance it needs to do the job well. At scale this becomes **context engineering** — designing ALL the info AI needs, not just wording one message. Memory aid: What -> How -> How to work with me.

Roman Urdu: Description = AI ko wo info aur guidance dena jo kaam achhe se karne ke liye chahiye. Scale par yeh context engineering ban jata hai.

- **Product description (EN):** What you want back — output type, audience, format, length, tone, topics to include/exclude.
  Roman Urdu: Jo output chahiye define karna — type, audience, format, length, tone.
- **Process description (EN):** How AI should approach the work — steps, order, method, checks before finishing.
  Roman Urdu: AI kaam kaise approach kare — steps, order, method, checks define karna.
- **Performance description (EN):** How AI should behave, for you and for its own users (concise/detailed, challenge vs agree, flag uncertainty).
  Roman Urdu: AI apne aap kaise behave kare, uske users ke liye — yeh define karna.

Completeness matters far more than clever wording. A good system prompt makes performance description permanent; a SKILL.md makes process description reusable.

3.3 Discernment — Judging the Output

EN: Discernment means judging the quality of what AI gave you. AI often sounds confident even when wrong — **automation bias** is trusting an automated answer too easily because it sounds professional.

Roman Urdu: Discernment = judge karna AI ne jo diya uski quality kaisi hai. Automation bias = confident/professional lagne wale answer ko bohot asaani se trust karna.

4 Signs of a Made-Up Answer (EN): (1) suspiciously exact specifics, (2) confidence where an expert would hesitate, (3) contradictions within a long output, (4) a claimed action that didn't actually happen.
Roman Urdu: (1) bohot exact specifics, (2) jahan expert hesitate kare wahan confidence, (3) long output mein contradiction, (4) claimed action jo hui hi nahi.

- **Product discernment (EN):** Is the result factually correct, complete, consistent, credible?
  Roman Urdu: Result khud achha/correct/consistent hai ya nahi, judge karna.
- **Process discernment (EN):** Is this way of working with AI actually paying off (vs. costing more time than it saves)?
  Roman Urdu: Kaam karne ka yeh tareeqa pay off kar raha hai ya nahi, judge karna.
- **Performance discernment (EN):** Does AI's independent behavior serve people well, across many cases (not just one output)?
  Roman Urdu: AI apne aap act karte hue log achhe se serve ho rahe hain ya nahi, judge karna.

Feedback pattern: **Problem -> Why it matters -> Direction** (not just “wrong, try again”).

3.4 Diligence — Owning the Use and Outcome

EN: Diligence asks a different question: “Should I use AI this way at all?” — taking responsibility for how AI is used and what happens with its output. Good output does NOT repair an irresponsible choice of tool or data.

Roman Urdu: Diligence = responsibility lena ke AI kaise use hua aur uske output ka kya hota hai. Achha output bhi galat tool/data choice ko theek nahi karta.

- **Creation diligence (EN):** Responsible choices about tools/data BEFORE and DURING use — is this data allowed in this tool? Use **redaction** (removing identifying details, keeping the task pattern) when possible.
  Roman Urdu: Pehle/during tools aur data responsibly choose karna; redaction = identifying details hata kar pattern rakhna.
- **Transparency diligence (EN):** Being honest about AI's role when it materially affects others (hiring, academic work, customer comms, medical/financial advice).
  Roman Urdu: Jab AI doosron ko affect kare, uske role ke baare mein honest hona.
- **Deployment diligence (EN):** Checking AI-assisted work before it's published/sent/executed — more impact or irreversibility = deeper check.
  Roman Urdu: Use/publish hone se pehle AI-assisted work verify karna aur ownership lena.

Numbers rule: any figure that's a decision-base (total/percentage) must be COMPUTED, never generated. **Litmus test:** “Would I confidently put my name on this?” — if no, it's not ready. **AI can automate work. It cannot automate accountability.**

4. Full Terms Cheatsheet (27 Terms)

One line each — English + Roman Urdu, straight from the course glossary.

Term	English	Roman Urdu
AI fluency	Working with AI effectively, efficiently, ethically, and safely.	AI ke sath effectively, efficiently, ethically, safely kaam karna.
The 4Ds	Delegation, Description, Discernment, Diligence.	Delegation, Description, Discernment, Diligence.
Automation	AI performs a task exactly as specified.	AI specify kiya task perform karta hai.
Augmentation	Human and AI work as thinking partners.	Human aur AI thinking partners ki tarah kaam karte hain.
Agency	AI works toward a goal, choosing many of its own steps.	AI goal ki taraf kaam karta hai, many steps khud choose karta hai.
Delegation	Deciding what AI should do, what humans keep.	Decide karna AI kya kare, humans kya rakhein.
Problem awareness	Knowing the goal, task, risks, and what success means.	Goal, kaam, risks, aur success ka matlab jaanna.
Platform awareness	Knowing which AI system/tool fits the task.	Kaunsa AI system/tool task ke liye fit hai, jaanna.
Task delegation	Assigning each part of work to human or AI.	Har part human ya AI ko assign karna.
Description	Giving AI the info and guidance it needs for the work.	AI ko info aur guidance dena jo kaam ke liye chahiye.
Product description	Defining the output you want.	Jo output chahiye, define karna.
Process description	Defining how AI should approach the work.	AI kaam kaise approach kare, define karna.
Performance description	Defining how AI should behave, for its users.	AI apne aap kaise behave kare, uske users ke liye define karna.
Discernment	Judging whether AI's output/justification/behavior is good enough.	Judge karna AI ka output/justification/behavior kaafi achha hai ya nahi.
Product discernment	Judging the result itself.	Result khud judge karna.
Process discernment	Judging whether working with AI this way is paying off.	Judge karna kya AI ke sath kaam karne ka tareeqa pay off kar raha.
Performance discernment	Judging whether AI's independent behavior serves users.	Judge karna AI ka independent behavior users ko serve karta hai ya nahi.
Diligence	Taking responsibility for how AI was used and its output.	AI kaise use hua aur uske output ki responsibility lena.
Creation diligence	Choosing tools and data responsibly before/during use.	Pehle/during tools aur data responsibly choose karna.
Transparency diligence	Being honest about AI's role when it affects others.	Jab affect kare, AI ke role ke baare mein honest hona.
Deployment diligence	Checking AI-assisted work before it's used/sent.	Use/send hone se pehle AI-assisted work check karna.
Context engineering	Designing all the info AI needs: instructions, docs, tools, memory, policies.	AI ko chahiye saari info design karna.
Automation bias	Trusting automated output too easily.	Insaan ka automated output ko bohot asaani se trust karna.
Hallucination	Confident AI output that sounds right but is invented/incorrect.	Confident AI output jo sound right karta par invented/incorrect hai.
Redaction	Removing identifying details before giving AI data, keeping the task pattern.	AI ko data dene se pehle identifying details hatana, task ka pattern rakh kar.
Digital FTE	An AI worker that does a defined job for others (FTE = full-time equivalent).	AI worker jo doosron ke liye defined job karta hai.
Tokens	Small text pieces AI reads/writes, which you pay for.	Chhote text pieces jo AI parhta/likhta hai, jinka payment hota hai.
System of Record	A business's trusted, official store of data and rules.	Business ka official data+rules ka trusted store.
System prompt	Standing instruction AI reads at the start of every chat.	Standing instruction jo AI har chat ki shuruat mein parhta hai.

Term	English	Roman Urdu
Eval suites	Repeatable tests that score an AI system's output.	Repeatable tests jo AI system ka output score karte hain.

5. Operating Loop, 10-80-10 Rule & Beginner Mistakes

The Loop (EN): Delegate -> Describe -> Discern -> Be diligent -> Repeat as needed. A personal skill becomes a system property when it scales into a Digital FTE.
Roman Urdu: Delegate -> Describe -> Discern -> Diligent raho -> zaroorat ho to repeat karo.

Chat vs. Agent Factory Mapping

Competency	In Chat	In Agent Factory
Delegation	Decide what to ask AI	Digital FTE scope + human/AI boundary
Description	Give instructions + context	System prompts, skills, context engineering, SoR
Discernment	Review the answer	Evals, monitoring, sampling, trusting the checker
Diligence	Protect data, own the result	Governance, permissions, audit, disclosure

10-80-10 Rule (maps to the Four Survival Skills)

- **First 10%** — set direction -> Delegation + Description are strongest.
- **Middle 80%** — orchestrate AI -> Description + Discernment repeat.
- **Last 10%** — judge the truth -> Discernment is critical.
- **Across 100%** — act responsibly -> **Diligence is NOT a final checkbox**, it surrounds the whole workflow (treating it as a final gate is Beginner Mistake #4).

4 Common Beginner Mistakes

Mistake	Missing Skill	Fix
Prompting before defining the problem	Delegation	Decide goal, audience, constraints, human/AI split first
Treating the first answer as final	Description+Discernment loop	Inspect result, give specific feedback, try again
Trusting a polished answer because it sounds professional	Discernment	Verify key facts, assumptions, calculations, sources
Only thinking about privacy/accountability after something goes wrong	Diligence	Decide data/disclosure/approval/accountability rules BEFORE deploying

Daily Checklist (make these automatic)

- **Delegate:** What's the goal? What should AI do? What stays with me?
- **Describe:** Output, context, method, behavior — what does AI need?
- **Discern:** Is the answer correct/complete/useful — how will I know?
- **Be diligent:** Is data safe? Should AI's role be disclosed? Who approves/owns this?

6. Viva Prep — 12 Scenario Q&A (Book's Own Quiz, Condensed)

Format: Question gist -> Correct answer -> One-line reasoning. Full verbatim options are in *09-test-your-understanding.md*.

#	Scenario Gist	Answer	One-line Reasoning
Q1	Manager asks AI to 'build a system for follow-ups'; output looks good but unusable.	**A — Delegation**	Who to contact, after how long, what tone, what needs approval are business policy questions prompting can't answer.
Q2	What separates Agency from Automation most sharply?	**D**	You're not in the room, and the work may be for others — judgment must be built in beforehand.
Q3	AI tutor's answers are all factually correct, yet students perform worse.	**C — Performance discernment**	Correct answers + worse outcomes = a behavior problem, visible only across many cases, not in any single output.
Q4	Writer spends 40 min steering AI on a 600-word article she could've drafted in 30.	**A — Process discernment**	Steering that costs more than it saves is a loss that feels like a win because it felt productive.
Q5	Telling AI 'don't just agree with me, flag uncertainty' is a scaled-down version of what?	**D — Performance description**	You're writing a behavior rule for the future; at scale the same sentence becomes an agent's standing rule.
Q6	AI's Vendor-A recommendation is sound, but based on a feature that doesn't exist.	**D**	A right answer resting on a wrong assumption won't stay right — judge the case made for it, not just the result.
Q7	Lecturer pastes student names/grades/notes into an unapproved consumer AI tool; feedback is excellent.	**A — Creation diligence**	Good output doesn't repair a bad choice of tool and data — the wrong data went into the wrong place.
Q8	Coding assistant introduces a security flaw; engineer ships it anyway.	**B**	'AI can automate work but not accountability' — the engineer and org still own the consequence.
Q9	Why is a new colleague like, and unlike, an AI assistant?	**C**	Both need context you forgot to give; the difference is memory — AI keeps nothing between chats, a colleague does.
Q10	Accountant reviews successful-looking matches too, not just failures, in a reconciliation agent test.	**B**	A system can fail silently — a repeated error can hide among the cases that look fine from the outside.
Q11	Comparing a vague prompt vs. a complete one (audience/time/topics/format stated).	**B**	The complete one isn't cleverer, it's more complete — completeness beats clever wording almost every time.
Q12	Where does Diligence sit in the 10-80-10 Rule?	**C — Across all 100%**	It surrounds the whole workflow rather than any one stage — treating it as a final checkbox is Beginner Mistake #4.

7. Quick Self-Check (10 Rapid-Fire Qs)

Question	Answer
AI fluency's 4 qualities?	Effective, efficient, ethical, safe.
Automation vs. augmentation vs. agency?	Automation = defined task; augmentation = thinking partner; agency = goal-directed, AI chooses steps, often for others.
Delegation's 3 parts?	Problem awareness, platform awareness, task delegation.
Description's 3 parts?	Product description, process description, performance description.
Why verify even a confident AI answer?	Sounding right != being correct — fluent wording doesn't verify facts/assumptions/reasoning.
Discernment's 3 parts?	Product discernment, process discernment, performance discernment.
Diligence's 3 parts?	Creation diligence, transparency diligence, deployment diligence.
Question to ask before shipping?	“Would I confidently put my name on this?”
The 4D loop in one line?	Decide what AI should do, describe the work, evaluate what comes back, own the whole process.
How does Discernment become engineering at scale?	It becomes eval suites, monitoring, sampling, and review gates that test whether the AI system performs well enough.

8. Where This Leads / Sources

Leads to: AI Prompting in 2026 (Description -> habits), How to Think in the AI Era (Discernment), Spec-Driven Development (Delegation+Description), Trusting the Checker + Eval-Driven Development (Discernment), System of Record + Designing the Vertical SoR (Diligence+governance).

Framework by: Prof. Rick Dakan (Ringling College) & Prof. Joseph Feller (University College Cork), with Anthropic. Free course: *AI Fluency: Framework and Foundations*, Claude Academy. Site: aifluencyframework.org. CC BY-NC-SA 4.0 / CC BY-NC-ND 4.0.
Generated from: docs/ai-fluency-crash-course/ (this repo) — content originally sourced via Zia Tutor AI connector, 2026-10-23.