Designing Agent Experiences

Mode 2 Phase 3 (Scale the Workforce), Chapter 2/9 — Revision & Viva-Prep Cheatsheet (English + Roman Urdu)

Source: agentfactory.panaversity.org/docs/designing-agent-experiences-crash-course

EN: An agentic product has 2 users at once: a human to earn trust from, and agents to parse correctly -- the surface must serve both without deceiving either. Trust = Reliability x Transparency x Control x Mistakes (a product, not a sum -- one zero ruins it).
Roman Urdu: Agentic product ke 2 users ek sath hote: insaan jise trust karna, aur agents jinhe parse karna -- surface dono ki khidmat kare bina kisi ko dhoka diye. Trust = Reliability x Transparency x Control x Mistakes (product hai, sum nahi -- ek zero poora zero banata).

1. The Shift + Human Surface Basics

- Third paradigm: Batch -> Command (you drove) -> Agentic (you state outcomes, agent chooses steps). 3 new design things: intent, trust, recovery.
- Autonomous agents create 3 NEW interfaces: Configuration, Monitoring, Intervention. **Progressive Transparency, 4 layers:** Outcome (1 line) -> Plan (steps) -> Why (honest confidence) -> Evidence (sources/trace).

2. Autonomy Dial + Intent Preview + Repair

- Autonomy is a DIAL, not a switch -- human-in-the-loop (pauses) vs human-on-the-loop (acts, human watches). High-stakes/irreversible always stays in-the-loop.
- Agent WILL make mistakes, not might. **4 repair moves:** undo (1-click, strongest trust-builder), apology, stated corrective action, visible human path.
- Supervising many workers shifts monitoring to TRIAGING: Fleet view, Attention triage, Drift signal (escalation rate doubling catches over-trust).

3. Machine Surface (AX) + Generative UI

- AX's 4 decisive factors: Access (scoped revocable credential), Context, Tools (typed/discoverable), Orchestration (chainable/idempotent). MCP deliberately leaves out orchestration, governance, state.
- MCP Apps give an interactive ui:// resource alongside tool text -- use widgets only for complex data/multi-option config/real-time monitoring, otherwise plain text is enough.

4. The New Craft + Measuring Experience

- New design objects: Policy surfaces, Confidence conveyors, System temperament -- you become a choreographer, not a screen-crafter.
- **Scorecard:** plan-acceptance rate, intervention rate, recovery success, over/under-trust, notification precision, time saved vs attention spent.
- **Anti-patterns:** black box, agentic sludge, over-eager agent, over-trusted agent, notification spam, trap door, confused deputy, mystery-meat API.

Key Terms Cheatsheet

One line each — English + Roman Urdu.

Term	English	Roman Urdu
Trust = Reliability x Transparency x Control x Mistakes	A product, not a sum -- one zero collapses the whole thing.	Product hai, sum nahi -- ek zero poora collapse kar deta.
Autonomy Dial	A spectrum, not a switch -- human-in-the-loop vs human-on-the-loop.	Spectrum hai, switch nahi -- human-in-the-loop vs human-on-the-loop.
AX (Agent Experience)	Design for the robot users of your product -- access, context, tools, orchestration.	Product ke robot users ke liye design -- access, context, tools, orchestration.
Agentic sludge	The cost of silently redistributing cognitive load -- everyone re-checks everything.	Cognitive load chupke se redistribute karne ka cost -- sab kuch sab dobara check karte.
Circuit breaker	Auto-lowers autonomy or pauses when a baseline-multiple is crossed.	Baseline-multiple cross hone par autonomy khud lower karta ya pause karta.

Quick Self-Check

Question	Answer
Why is trust modeled as a PRODUCT, not a sum, of its 4 factors?	One factor at zero (e.g. no recovery path) collapses overall trust, even if the other three are high.
What's the difference between human-in-the-loop and human-on-the-loop?	In-the-loop pauses for approval before acting; on-the-loop acts and lets the human watch and intervene -- high-stakes work always stays in-the-loop.
Why is 'agentic sludge' dangerous even though nothing visibly breaks?	Cognitive load got silently redistributed -- everyone ends up re-checking everything, quietly destroying the efficiency gain.
What does a doubled escalation rate signal in a fleet-monitoring surface?	Over-trust being caught -- a drift signal that autonomy was extended faster than the work actually supports.

Where This Leads / Sources

Leads to: Eval-Driven Development (correctness) vs this course (relationship/experience metrics). Full worked Support Worker example + Hands-On Lab in Part 4.

Source: docs/designing-agent-experiences-crash-course/ (this repo), content originally via Zia Tutor AI connector.